

# perf

(Linux profiling with performance counter)

---

Concurrent Programming

# Introduction

---

- What is Perf?
- Installing Perf
- Perf Commands
- Example

# What is Perf?

- Performance analyzing tool in Linux
  - available from Linux kernel version 2.6.31
- Capable of statistical profiling of the entire system
  - (both kernel and userland code)
- Available to measure many types of event
  - (Hardware events, Software events, ...)

# Perf Commands

- perf list  
Show the list of available events to be measured with perf
- perf stat  
Obtain aggregated event counts
- perf record  
Sampling events on per-thread, per-process, per-cpu basis and generate the output file
- perf report  
Analyze the output file generated from *perf record*
- perf annotate  
Analyze the output file generated from *perf record* in the instruction level
- more Commands, but not today

# Perf Commands – perf list

- Show the list of available events to be measured with perf

```
$ sudo perf list
```

```
mrbin2002@ubuntu:~/TA/Multicore$ sudo perf list

List of pre-defined events (to be used in -e):

alignment-faults          [Software event]
bpf-output                 [Software event]
context-switches OR cs    [Software event]
cpu-clock                 [Software event]
cpu-migrations OR migrations [Software event]
dummy                     [Software event]
emulation-faults          [Software event]
major-faults              [Software event]
minor-faults              [Software event]
page-faults OR faults     [Software event]
task-clock                [Software event]

msr/smi/                  [Kernel PMU event]
msr/tsc/                  [Kernel PMU event]
power/energy-cores/       [Kernel PMU event]
power/energy-gpu/         [Kernel PMU event]
power/energy-pkg/         [Kernel PMU event]

rNNN                      [Raw hardware event descriptor]
cpu/t1=v1[,t2=v2,t3 ...]/modifier [Raw hardware event descriptor]
(see 'man perf-list' on how to encode it)

mem:<addr>[/len][:access] [Hardware breakpoint]
```

# Perf Commands – perf stat

- Generate the statistics of the events that are occurred during process execution

```
$ sudo perf stat ./prac_mutex
```

```
mrbin2002@ubuntu:~/TA/Multicore/lab2$ sudo perf stat ./prac_mutex
thread 140504191280896, local count: 1000000
thread 140504182888192, local count: 1000000
thread 140504174495488, local count: 1000000
thread 140504166102784, local count: 1000000
thread 140504157710080, local count: 1000000
thread 140504149317376, local count: 1000000
thread 140504140924672, local count: 1000000
thread 140504132531968, local count: 1000000
thread 140504124139264, local count: 1000000
thread 140504115746560, local count: 1000000
global count: 10000000

Performance counter stats for './prac_mutex':

      3303.401893    task-clock (msec)    #    3.393 CPUs utilized
           303,808    context-switches    #    0.092 M/sec
          196,931    cpu-migrations     #    0.060 M/sec
              82     page-faults      #    0.025 K/sec
<not supported>    cycles
<not supported>    stalled-cycles-frontend
<not supported>    stalled-cycles-backend
<not supported>    instructions
<not supported>    branches
<not supported>    branch-misses

      0.973515568 seconds time elapsed
```

# Perf Commands – perf record

- Sampling events on per-thread, per-process, per-cpu basis and generate the output file

```
$ sudo perf record -g ./prac_mutex
```

```
[mrbin2002@ubuntu:~/TA/Multicore/lab2$ sudo perf record -g ./prac_mutex
thread 139750672590592, local count: 1000000
thread 139750664197888, local count: 1000000
thread 139750655805184, local count: 1000000
thread 139750647412480, local count: 1000000
thread 139750639019776, local count: 1000000
thread 139750630627072, local count: 1000000
thread 139750622234368, local count: 1000000
thread 139750613841664, local count: 1000000
thread 139750605448960, local count: 1000000
thread 139750597056256, local count: 1000000
global count: 10000000
[ perf record: Woken up 5 times to write data ]
[ perf record: Captured and wrote 1.336 MB perf.data (13013 samples) ]
[mrbin2002@ubuntu:~/TA/Multicore/lab2$ ls
Makefile perf.data prac_mutex prac_mutex.cpp prac_mutex.o
```

# Perf Commands – perf record

- Sampling events on a running process

```
$ sudo perf record -g -p 7553
```

```
[mrbin2002@ubuntu:~/TA/Multicore/lab2$ ps u
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
mrbin20+  2491  0.0  0.0  28056  6716 pts/1    Ss   15:15   0:00 -bash
mrbin20+  7425  0.2  0.0  28056  6720 pts/10   Ss   15:39   0:00 -bash
mrbin20+  7553  403  0.1  94812  9656 pts/1    Sl+  15:40   0:16 ./prac_mutex
mrbin20+  7576  0.0  0.1  30904  7160 pts/10   R+   15:40   0:00 ps u
[mrbin2002@ubuntu:~/TA/Multicore/lab2$ sudo perf record -g -p 7553
[ perf record: Woken up 17 times to write data ]
[ perf record: Captured and wrote 2.006 MB perf.data (19992 samples) ]
```

```
$ sudo perf record -g -p `pidof prac_mutex`
```



# Perf Commands – perf report

- Analyze the output file generated from perf record

```
$ sudo perf report -g graph --no-children
```

```
Samples: 13K of event 'cpu-clock', Event count (approx.): 3288500000
```

Overhead	Command	Shared Object	Symbol
+ 16.32%	prac_mutex	[kernel.kallsyms]	[k] _raw_spin_unlock_irqrestore
+ 13.25%	prac_mutex	libpthread-2.19.so	[.] pthread_mutex_lock
+ 9.21%	prac_mutex	libpthread-2.19.so	[.] pthread_mutex_unlock
+ 9.05%	prac_mutex	[kernel.kallsyms]	[k] native_queued_spin_lock_slowpath
+ 6.53%	prac_mutex	libpthread-2.19.so	[.] __l1ll_lock_wait
+ 6.01%	prac_mutex	[kernel.kallsyms]	[k] futex_wake
+ 5.99%	prac_mutex	[kernel.kallsyms]	[k] finish_task_switch
+ 5.55%	prac_mutex	[kernel.kallsyms]	[k] futex_wait_setup
+ 4.85%	prac_mutex	[kernel.kallsyms]	[k] _raw_spin_lock
+ 4.09%	prac_mutex	[kernel.kallsyms]	[k] entry_SYSCALL_64_after_swapgs
+ 2.97%	prac_mutex	[kernel.kallsyms]	[k] get_futex_key_refs.isra.13
+ 2.73%	prac_mutex	[kernel.kallsyms]	[k] hash_futex
+ 1.96%	prac_mutex	[kernel.kallsyms]	[k] get_futex_key
+ 1.96%	prac_mutex	[kernel.kallsyms]	[k] get_futex_value_locked
+ 1.47%	prac_mutex	libpthread-2.19.so	[.] __l1ll_unlock_wake
+ 1.34%	prac_mutex	[kernel.kallsyms]	[k] futex_wait
+ 1.34%	prac_mutex	prac_mutex	[.] _Z10ThreadFuncPv

# Perf Commands – perf annotate

- Analyze the output file generated from perf record in the instruction level

```
$ sudo perf annotate pthread_mutex_lock
```

```
pthread_mutex_lock /lib/x86_64-linux-gnu/libpthread-2.19.so
```

```
      lea    0x16(%r8),%rsi
      mov    %r8,%rdi
      and    $0x80,%edx
      add    $0x8,%rsp
      ↓ jmpq  7af0
      nop
58:    and    $0x80,%esi
0.40   mov    $0x1,%edi
0.40   xor    %eax,%eax
      lock  cmpxchg %edi,(%r8)
59.04 ↓ jne    23a
0.40   mov    0x8(%r8),%eax
9.24   test   %eax,%eax
      ↓ jne    dd
0.40   78:    mov    %fs:0x2d0,%eax
0.29   addl   $0x1,0xc(%r8)
1.38   mov    %eax,0x8(%r8)
0.57   xor    %eax,%eax
      8b:    add    $0x8,%rsp
      ← retq
90:    cmp    $0x100,%eax
```

# Perf Commands – perf annotate

- Analyze the output file generated from perf record in the instruction level

You can run the *annotate* command simply in the report screen

Samples: 13K of event 'cpu-clock', Event count (approx.): 3288500000

Overhead	Command	Shared Object	Symbol
+ 16.32%	prac_mutex	[kernel.kallsyms]	[k] _raw_spin_unlock_irqrestore
- 13.25%	prac_mutex	libpthread-2.19.so	[.] pthread_mutex_lock
pthread_mutex_lock			
start_thread			
+ 9.21%	prac_mutex	libpthread-2.19.so	[.] pthread_mutex_unlock
+ 9.05%	prac_mutex	[kernel.kallsyms]	[k] native_queued_spin_lock_slowpath
+ 6.53%	prac_mutex	libpthread-2.19.so	[.] __l1ll_lock_wait
+ 6.01%	prac_mutex	[kernel.kallsyms]	[k] futex_wake
+ 5.99%	prac_mutex	[kernel.kallsyms]	[k] finish_task_switch
+ 5.55%	prac_mutex	[kernel.kallsyms]	[k] futex_wait_setup
+ 4.85%	prac_mutex	[kernel.kallsyms]	[k] _raw_spin_lock
+ 4.09%	prac_mutex	[kernel.kallsyms]	[k] entry_SYSCALL_64_after_swapgs
+ 2.97%	prac_mutex	[kernel.kallsyms]	[k] get_futex_key_refs.isra.13
+ 2.73%	prac_mutex	[kernel.kallsyms]	[k] hash_futex
+ 1.96%	prac_mutex	[kernel.kallsyms]	[k] get_futex_key
+ 1.96%	prac_mutex	[kernel.kallsyms]	[k] get_futex_value_locked
+ 1.47%	prac_mutex	libpthread-2.19.so	[.] __l1ll_unlock_wake
+ 1.34%	prac_mutex	[kernel.kallsyms]	[k] futex_wait
+ 1.34%	prac_mutex	prac_mutex	[.] _Z10ThreadFuncPv
+ 0.97%	prac_mutex	[kernel.kallsyms]	[k] sys_futex



pthread_mutex_lock		/lib/x86_64-linux-gnu/libpthread-2.19.so	
		lea	0x16(%r8),%rsi
		mov	%r8,%rdi
		and	\$0x80,%edx
		add	\$0x8,%rsp
		↓ jmpq	7af0
		nop	
	58:	and	\$0x80,%esi
0.40		mov	\$0x1,%edi
0.40		xor	%eax,%eax
		lock	cmpxchg %edi, (%r8)
59.04		↓ jne	23a
0.40		mov	0x8(%r8),%eax
9.24		test	%eax,%eax
		↓ jne	dd
0.40	78:	mov	%fs:0x2d0,%eax
0.29		addl	\$0x1,0xc(%r8)
1.38		mov	%eax,0x8(%r8)
0.57		xor	%eax,%eax
	8b:	add	\$0x8,%rsp
		← retq	
	90:	cmp	\$0x100,%eax

Choose a symbol and press 'a' key

# More about Perf...

---

Perf Tutorial

<https://perf.wiki.kernel.org/index.php/Tutorial>

# Thank You

---